

## Basic Idea

This project is basically a personal life management tool with which the user will interact in natural language.

For example:

User says: I need to fix a bug in auth in avenue

Secretary: Created tasks “Auth bug in avenue”

User asks: What are all the pending tasks?

Secretary: Gives a list of all pending tasks

User says: I have an event on 16th October

Secretary: Saved that event

User asks: Am I free on 16th October

Secretary: No, you have an event on 16th October

Basically, this agent will store all the context of tasks and events

What exactly is happening?

1. **Conversation Engine (AI brain):** Understands what you said and decides *what kind of task* it is — e.g., query, creation, update, summary, etc.
2. **Memory + Logic Layer (Backend brain):** Actually *does* the work — saves tasks, fetches events, retrieves memory, etc.
3. **Data Storage (Long-term brain):** Where all your facts live — tasks, events, notes, embeddings (memory).

Logical Flow

1. You send a message (frontend)

Example:

“Remind me to fix the login bug in Product 1 by Friday.”

## Frontend (Next.js)

- Sends { message: "Remind me..." } → to your backend API (/api/chat).
- Shows your message in the chat UI.
- Waits for the assistant's response (streaming or full).

2. Backend receives it and creates a “conversation state”

## What the backend does:

- It collects recent chat history (maybe last 10 messages).
- Optionally fetches any relevant user memory (tasks/events).
- Combines that into a “prompt context”.

## Why?

So the model has enough information to understand *your* current world.

### 3. Backend sends everything to the LLM (like GPT-4)

It sends a request that looks roughly like this:

```
[  
  { role: "system", content: "You are a helpful secretary..." },  
  { role: "user", content: "Remind me to fix the login bug in Product 1 by Friday." }  
]
```

### 4. The LLM interprets the intent

Now the model decides:

→ Is this a **question**, **task creation**, **status update**, or **schedule query**?

Depending on that, it may:

- Ask for clarification (if it's ambiguous), **or**
- Decide to **call a tool/function** (e.g. `create_task`).

### 5. If it calls a tool, backend executes it

Example tool call:

```
{  
  "name": "create_task",  
  "arguments": {  
    "title": "Fix login bug in Product 1",  
    "dueAt": "2025-10-10T17:00:00Z"  
  }  
}
```

**Backend executes that:**

- Inserts into your tasks table.
- Generates an **embedding** (vector) for that text → stores it in memory.
- Returns the new task info to the LLM.

### 6. LLM receives the result and replies naturally

It now formats the response:

“Got it. I've added ‘Fix login bug in Product 1’ due Friday.”

Backend sends that back to frontend → chat updates.

## 7. Future queries use “memory”

Later you ask:

“What’s pending this week?”

Backend flow:

1. Detect this is a **status query**.
2. Calls database: find tasks where status != done and due\_at < end\_of\_week.
3. Optionally uses **vector search** for semantic matches (e.g., “finish product stuff”).

Sends retrieved items to the LLM:

```
{  
  role: "system": "...",  
  role: "user": "What's pending this week?",  
  role: "assistant": "context: [task1, task2, task3]"  
}
```

- 4.
5. LLM summarizes naturally:  
“You have 2 tasks due: fix the login bug in Product 1, and finalize the pricing doc for Project 2.”

## 8. Daily summary or proactive actions

You can have a cron job (or server trigger) that:

- Runs every morning.
- Queries DB for today’s schedule and overdue tasks.
- Uses the LLM to summarize:  
“Good morning! You have 2 meetings today and 1 overdue task from yesterday.”
- Sends that via chat, email, or notification.

## Putting It All Together — Component View

### 1. Frontend (Chat UI / Task View)

- Sends user messages to /api/chat
- Displays messages from assistant
- Shows structured data (tasks, calendar, etc.)

### 2. Backend API (Logic Layer)

- Routes chat messages
- Calls OpenAI / LLM APIs
- Handles tool calls (create/list/update)
- Manages DB operations

- Stores embeddings for memory

### 3. Database (Memory)

- Stores structured data (tasks, events)
- Stores vector embeddings for memory search
- Keeps chat history (optional)

### 4. LLM (Reasoning Engine)

- Interprets intent
- Decides when to call backend tools
- Generates human-friendly responses

## Pricing

### Stage 2: Fine-tuning (one-time, \$10–200)

OpenAI GPT-4o mini fine-tuning (updated 2026 rates): [openai](#)

Component	Price per 1M tokens	1M token dataset	10M token dataset
Training	\$5.00 <a href="#">openai</a>	\$5	\$50
Input inference (post-tune)	\$0.80 <a href="#">openai</a>	\$0.80	\$8
Output inference (post-tune)	\$3.20 <a href="#">openai</a>	\$3.20	\$32
Total	-	~\$9	~\$90

- Storage: Negligible (~\$2/month per model). [abot](#)
- Mistral Nemo alternative: \$1/M training (cheaper for small datasets). [abot](#)
- Scale to 100M tokens only if needed (rare for secretary tasks). [openai](#)

### **Light usage (MVP / side project)**

Say **100 interactions/day** (onboarding + some emails + FAQs):

- Daily tokens  $\approx 40,000$
- Monthly tokens  $\approx 1.2\text{M}$  total

Monthly cost:

- Input:  $0.15 \times 0.24 \approx 0.036$  USD
- Output:  $0.60 \times 0.24 \approx 0.144$  USD
- **Total  $\approx 0.18$  USD/month** (around ₹15). [\[1\]](#) [\[2\]](#)

### **Moderate usage (small team / startup pilot)**

Say **1,000 interactions/day**:

- Daily tokens  $\approx 400,000$
- Monthly tokens  $\approx 12\text{M}$
- Input:  $0.15 \times 2.4 \approx 0.36$  USD
- Output:  $0.60 \times 2.4 \approx 1.44$  USD
- **Total  $\approx 1.80$  USD/month** ( $\approx$  ₹150). [\[1\]](#) [\[2\]](#)

### **Heavy usage (10,000 interactions/day)**

Still very cheap:

- Tokens/month  $\approx 120\text{M}$
- Input:  $0.15 \times 24 \approx 3.60$  USD
- Output:  $0.60 \times 24 \approx 14.40$  USD